Search for articles...

All Collections > Consumers > Data Handling & Retention >
How Do You Use Personal Data in Model Training?

How Do You Use Personal Data in Model Training?

Updated this week

Table of contents

This article is about our consumer products such as Claude Free, Pro, Max and when accounts from those plans use Claude Code. For our commercial products such as Claude for Work and the Anthropic API, see here.

About Model Training

Large language models such as Claude are “trained” on a variety of content such as text, images and multimedia so that they can learn the patterns and connections between words and/or content. This training is important so that the model performs effectively and safely.

Models do not store text like a database, nor do they simply “mash-up” or “collage” existing content. Models identify general patterns in text in order to help people create new content, and they do not have access to or pull from the original training data once the models have been trained.

Collection of Personal Data

The following sources of training data may contain personal data:

1. Publicly available information via the Internet
2. Datasets that we obtain under commercial agreements with third party businesses
3. Data that our users or crowd workers provide
4. Data you allow us to use, such as your chats and coding sessions, to improve Claude.

We do not actively set out to use personal data to train our models. However, a large amount of data on the Internet relates to people, so our training data may incidentally include personal data. We only use personal data included in our training data to help our models learn about language and how to understand and respond to it. We do not use such personal data to contact people, build profiles about them, to try to sell or market anything to them, or to sell the information itself to any third party.

Privacy Safeguards During Data Collection and Training

We take steps to minimize the privacy impact on individuals through the training process. We operate under strict policies and guidelines, for instance, that we do not access password protected pages or bypass CAPTCHA controls. We undertake due diligence on the data that we license. And we encourage our users not to use our products and services to process personal data.

Additionally, our models are specifically trained to respect privacy. We have built key 'privacy by design' safeguards into the development of Claude through our adoption of "Constitutional AI". This gives Claude a set of principles (i.e., a "constitution") to guide the training of the Claude LLMs and to make judgments about outputs. These principles are based in part on the Universal Declaration of Human Rights and include specific rules around protecting privacy, particularly of non-public figures. This trains the Claude LLMs to not disclose or repeat personal data which may have been incidentally captured in training data, even if prompted. For example, Claude is given the following principles as part of its "constitution": "Please choose the response that is most respectful of everyone's privacy" and "Please choose the response that has the least personal, private, or confidential information belonging to others". For more information on how "Constitutional AI" works, see here.

Where you have allowed us to use your chats and coding sessions to improve Claude, we will automatically de-link them from your user ID (e.g. email address) before it's used by Anthropic.

We will use tools and processes derived from our work on privacy-preserving analysis tools to filter or obfuscate sensitive data. In addition to that, our pioneering research in post-training techniques helps to minimize the amount of personal data included in Claude's outputs.

Your data is used solely to make Claude better for everyone - we do not use such personal data to contact people, build profiles about them, to try to sell or market anything to them, or to sell the information itself to any third party.

Additionally, you maintain full control and can adjust your privacy settings at any time.

Data usage for Claude.ai Consumer Offerings (e.g. Claude, Pro, Max, etc.)

We may use your chats or coding sessions to improve our models, if:

1. You choose to allow us to use your chats and coding sessions to improve Claude,
2. Your conversations are flagged for safety review (in which case we may use or analyze them to improve our ability to detect and enforce our Usage Policy, including training models for use by our Safeguards team, consistent with Anthropic's safety mission),
3. You've explicitly provided materials to us (e.g.via our thumbs up/down feedback button), or
4. By otherwise explicitly opting in to training (e.g. by joining our Trusted Tester Program).

Privacy Rights and Data Processing

Our Privacy Policy explains your rights regarding your personal data, including with respect to our training activities. This includes your right to request a copy of your personal data, and to object to our processing of your personal data or request that it is deleted. We make every effort to respond to such requests. However, please be aware that these rights are limited, and that the process by which we may need to action your requests regarding our training dataset are complex.

To learn more, or if you would like to know how to contact us regarding a privacy related topic, see our Trust Center and Privacy Policy.

Please note, the Privacy Policy does not apply where Anthropic acts as a data processor and processes personal data on behalf of commercial customers using Anthropic's Commercial Services. In those cases, the commercial customer is the controller, and you can review their policies for more information about how they handle your personal data.

Related Articles

- Is my data used for model training?
- How do you use personal data in model training?
- How long do you store my data?
- Is my data used for model training?
- How do I change my model improvement privacy settings?

Did this answer your question?

😞 😐 😃